

R-Analysis

Tom Blain, Daniel Gardner, Xin Guan, Xinyu Li

05/10/2022

```
#if(!require("readr")) install.packages("readr")
#if(!require("ggplot2")) install.packages("ggplot2")
#if(!require("gridExtra")) install.packages("gridExtra")
#if(!require("dplyr")) install.packages("dplyr")
#if(!require("ggplots")) install.packages("ggplots")
#if(!require("plotrix")) install.packages("plotrix")
library("fs") # for cross-platform directories (path_wd)
library("readr") #For read_csv
library("dplyr")
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library("knitr") # For kable
library("ggplot2") # For plots
library("gridExtra")
```

```
##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##   combine
```

```
library("gplots")
```

```
##
## Attaching package: 'gplots'

## The following object is masked from 'package:stats':
##
##   lowess
```

```
library("plotrix") # For general stacked histogram
```

```
##
## Attaching package: 'plotrix'

## The following object is masked from 'package:gplots':
```

```
##
##      plotCI
data <- read_csv(path_wd("01-Data.csv"))

## Rows: 891 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr (5): Name, Sex, Ticket, Cabin, Embarked
## dbl (7): PassengerId, Survived, Pclass, Age, SibSp, Parch, Fare
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
data<-as.data.frame(data)
```

Analysing the Data

Let's start with looking at first few rows of data.

```
head(data)
```

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch
1	0	3	Braund, Mr. Owen Harris	male	22	1	0
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38	1	0
3	1	3	Heikkinen, Miss. Laina	female	26	0	0
4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0
5	0	3	Allen, Mr. William Henry	male	35	0	0
6	0	3	Moran, Mr. James	male	NA	0	0

Ticket	Fare	Cabin	Embarked
A/5 21171	7.2500	<NA>	S
PC 17599	71.2833	C85	C
STON/O2. 3101282	7.9250	<NA>	S
113803	53.1000	C123	S
373450	8.0500	<NA>	S
330877	8.4583	<NA>	Q

Here is a brief summary of the data set.

```
summary(data)
```

PassengerId	Survived	Pclass	Name
Min. : 1.0	Min. :0.0000	Min. :1.000	Length:891
1st Qu.:223.5	1st Qu.:0.0000	1st Qu.:2.000	Class :character
Median :446.0	Median :0.0000	Median :3.000	Mode :character
Mean :446.0	Mean :0.3838	Mean :2.309	
3rd Qu.:668.5	3rd Qu.:1.0000	3rd Qu.:3.000	
Max. :891.0	Max. :1.0000	Max. :3.000	

Sex	Age	SibSp	Parch
-----	-----	-------	-------

```
## Length:891      Min.   : 0.42   Min.   :0.000   Min.   :0.0000
## Class :character 1st Qu.:20.12   1st Qu.:0.000   1st Qu.:0.0000
## Mode :character Median :28.00   Median :0.000   Median :0.0000
##                Mean  :29.70   Mean  :0.523   Mean  :0.3816
##                3rd Qu.:38.00   3rd Qu.:1.000   3rd Qu.:0.0000
##                Max.   :80.00   Max.   :8.000   Max.   :6.0000
##                NA's   :177
## Ticket          Fare          Cabin          Embarked
## Length:891      Min.   : 0.00   Length:891      Length:891
## Class :character 1st Qu.: 7.91   Class :character Class :character
## Mode :character Median :14.45   Mode :character Mode :character
##                Mean  :32.20
##                3rd Qu.:31.00
##                Max.   :512.33
##
```

This data set consists of our binary survival variable we are interested, as well as 9 other co-variates that may influence survival. We can start by encoding our categorical variables as factors, with the ticket class variable being ordered in this case.

```
data$Sex <- as.factor(data$Sex)
data$Survived <- as.factor(data$Survived)
data$Pclass <- as.ordered(data$Pclass)
```

Then we can check the incompleteness of the data,

```
apply(data, function(x) {sum(is.na(x))})
```

```
## PassengerId    Survived    Pclass      Name      Sex      Age
##           0           0           0           0           0      177
##      SibSp      Parch      Ticket      Fare      Cabin      Embarked
##           0           0           0           0      687           2
```

We can see that for the most part this data set is complete. Importantly, key variables like Survived, Pclass and Sex have all been recorded, besides the 891 NA observations we are using for the test data. However there are quite large gaps when it comes to Age (Missing 177) and Cabin number (Missing 687).

We can deal with missing values in each column by fill a suitable substitution such as mode, mean or median.

```
data$Embarked[is.na(data$Embarked)]<-mode(data$Embarked)
#Use mode for the missing "Embarked" values.

data$Age[is.na(data$Age)]<-mean(data$Age,na.rm = T)
#Use average of the existing age values for the missing "Age" values.
```

Manipulating the data

We can also try to deconstruct some of the variables to infer more information. For example, the passenger name is very difficult to use in any kind of analysis, so we can try and turn it into a factor variable based off the title of each person.

```
data$Title <- sapply(data$Name, function(x) {strsplit(x, split='[,.]')[[1]][2]})
data$Title <- sub(' ', '', data$Title) #removing spaces before title
kable(table(data$Sex, data$Title))
```

	Capt	Col	Don	Dr	Jonkheer	Lady	Major	Master	Miss	Mlle	Mme	Mr	Mrs	Ms	Rev	Sir
female	0	0	0	1	0	1	0	0	182	2	1	0	125	1	0	0

	Capt	Col	Don	Dr	Jonkheer	Lady	Major	Master	Miss	Mlle	Mme	Mr	Mrs	Ms	Rev	Sir
male	1	2	1	6	1	0	2	40	0	0	0	517	0	0	6	1


This is better than having over 1700 different names, although we can still reduce the number of dimensions of this new title variable. One way of doing this is grouping similar titles together, such as ‘Ms’, ‘Miss’ and ‘Mlle’ (Mademoiselle), all referring to young or unmarried women. Similarly we can group ‘Mrs’ and ‘Mme’ (Madame) as these refer to married women. Anything beside these two categories along with Mr and Master we can group in one ‘Rare titles’ category as miscellaneous titles.

```
data$Title[data$Title %in% c("Mlle", "Ms")] <- "Miss"
data$Title[data$Title == "Mme"] <- "Mrs"
data$Title[!(data$Title %in% c('Master', 'Miss', 'Mr', 'Mrs'))] <- "Rare Title"
data$Title <- as.factor(data$Title)
kable(table(data$Sex, data$Title))
```

	Master	Miss	Mr	Mrs	Rare Title
female	0	185	0	126	3
male	40	0	517	0	20

Now we have a much more clear table we can plot this to get an idea of how Title affects survival

```
ggplot(data[!is.na(data$Survived),], aes(x = Title, fill = Survived)) +
  geom_bar(stat='count', position='stack') +
  labs(x = 'Title') + theme_grey()
```



Create a variable “Famillysize” which is the sum of variables “SibSp” and “Parch”. Analysing family size

seems to make more sense than analyse siblings and parents separately.


```
data$Familysize=data$SibSp+data$Parch
```

Visualising the data

We can then start visualizing the amount of deaths and who died due to the factor variables via bar charts

Let's firstly look at the most basic stacked histogram: We see that the survival rate of female is a lot higher than male, this is because that the men on board the ocean liner gave women and children priority access to the lifeboats.


```
histStack(x=data$Fare,z=factor(data$Survived),col=c("red","blue"),main = "Fare Histogram by Survived",ylim=c(0,800),  
legend("topright",legend=c("Survived","Dead"),col=c("blue","red"),lwd=5,cex=0.5) #Add legend to indicate survival status
```



People between 20-40 years old has significantly lower survival rate.


```
histStack(x=data$Age,z=factor(data$Survived),col=c("red","blue"),main = "Age Histogram by Survived",ylim=c(0,800),  
legend("topright",legend=c("Survived","Dead"),col=c("blue","red"),lwd=5,cex=0.5)
```

Age Histogram by Survived



```
histStack(x=data$Familysize,z=factor(data$Survived),col=c("red","blue"),main = "Family size Histogram by Survived",
legend("topright",legend=c("Survived", "Dead"),col=c("blue", "red"),lwd=5,cex=0.5)
```


Family size Histogram by Survived



If we zoom in a little bit, we can see that the family size of 2 or 3 has significantly higher survival rate than family of other sizes.


```
histStack(x=data$Familysize,z=factor(data$Survived),col=c("red","blue"),main = "Family size Histogram by Survived",
legend("topright",legend=c("Survived", "Dead"),col=c("blue", "red"),lwd=5,cex=0.5)
```

Family size Histogram by Survived



There are other options such as `ggplot2`, which plots prettier graphs, and contains more graphical features.

```
ggplot(data[!is.na(data$Survived),], aes(x = Survived, fill = Survived)) +  
  geom_bar(stat='count') +  
  labs(x = 'How many people died and survived on the Titanic?') +  
  geom_label(stat='count', aes(label=..count..), size=7) +  
  theme_grey(base_size = 18)
```



How many people died and survived on the Titanic?

And we can do further analysis on how class affects your chances of survival

```
p3 <- ggplot(data, aes(x = Pclass, fill = Pclass)) +
  geom_bar(stat='count', position='dodge') +
  labs(x = 'Pclass') + geom_label(stat='count', aes(label=..count..)) +
  theme(legend.position="none") + theme_grey()
p4 <- ggplot(data[!is.na(data$Survived),], aes(x = Pclass, fill = Survived)) +
  geom_bar(stat='count', position='dodge') + labs(x = 'Pclass') +
  theme(legend.position="none") + theme_grey()
p5 <- ggplot(data[!is.na(data$Survived),], aes(x = Pclass, fill = Survived)) +
  geom_bar(stat='count', position='stack') +
  labs(x = 'Pclass', y = "Count") + facet_grid(.~Sex) +
  theme(legend.position="none") + theme_grey()
p6 <- ggplot(data[!is.na(data$Survived),], aes(x = Pclass, fill = Survived)) +
  geom_bar(stat='count', position='fill') +
  labs(x = 'Pclass', y = "Percent") + facet_grid(.~Sex) +
  theme(legend.position="none") + theme_grey()

grid.arrange(p3, p4, p5, p6, ncol=2)
```